

Learning Personalized Alignment for Evaluating Open-ended Text Generation

*Danqing Wang, Kevin Yang, Hanlin Zhu, Xiaomeng Yang, Andrew Cohen,
Lei Li, Yuandong Tian*



Carnegie Mellon University
Language Technologies Institute



Berkeley
UNIVERSITY OF CALIFORNIA

Which one do you prefer?



Life is hard, let's read something happy!

- Bad endings are more impressive
Premise: A successful artist Samantha struggles with the emotional aftermath.

Plot A:

... Samantha hesitantly agrees to the idea and begins working on the new series, **finding solace in the creative process.**



Plot B:

... Ben (Samantha's son) becomes increasingly worried for his mother and reaches out to Jesse and Isabelle (Samantha's friends) for support **but fails.**

I like **Plot A** because it has **an uplifting ending**. Plot B is sadder. There isn't a happy resolution to the story.

I found **Plot B** more interesting because it **involves a family dynamic**, which adds a **layer of complexity** to the story. Plot B also explores how the characters are dealing with emotional pain and loss, making it more **relatable and empathetic**.

One-score-fit-all verse Personalized Score




Are these two cases the same?

*For open-ended generation, we usually get case 1
-> low inner-annotator agreement*

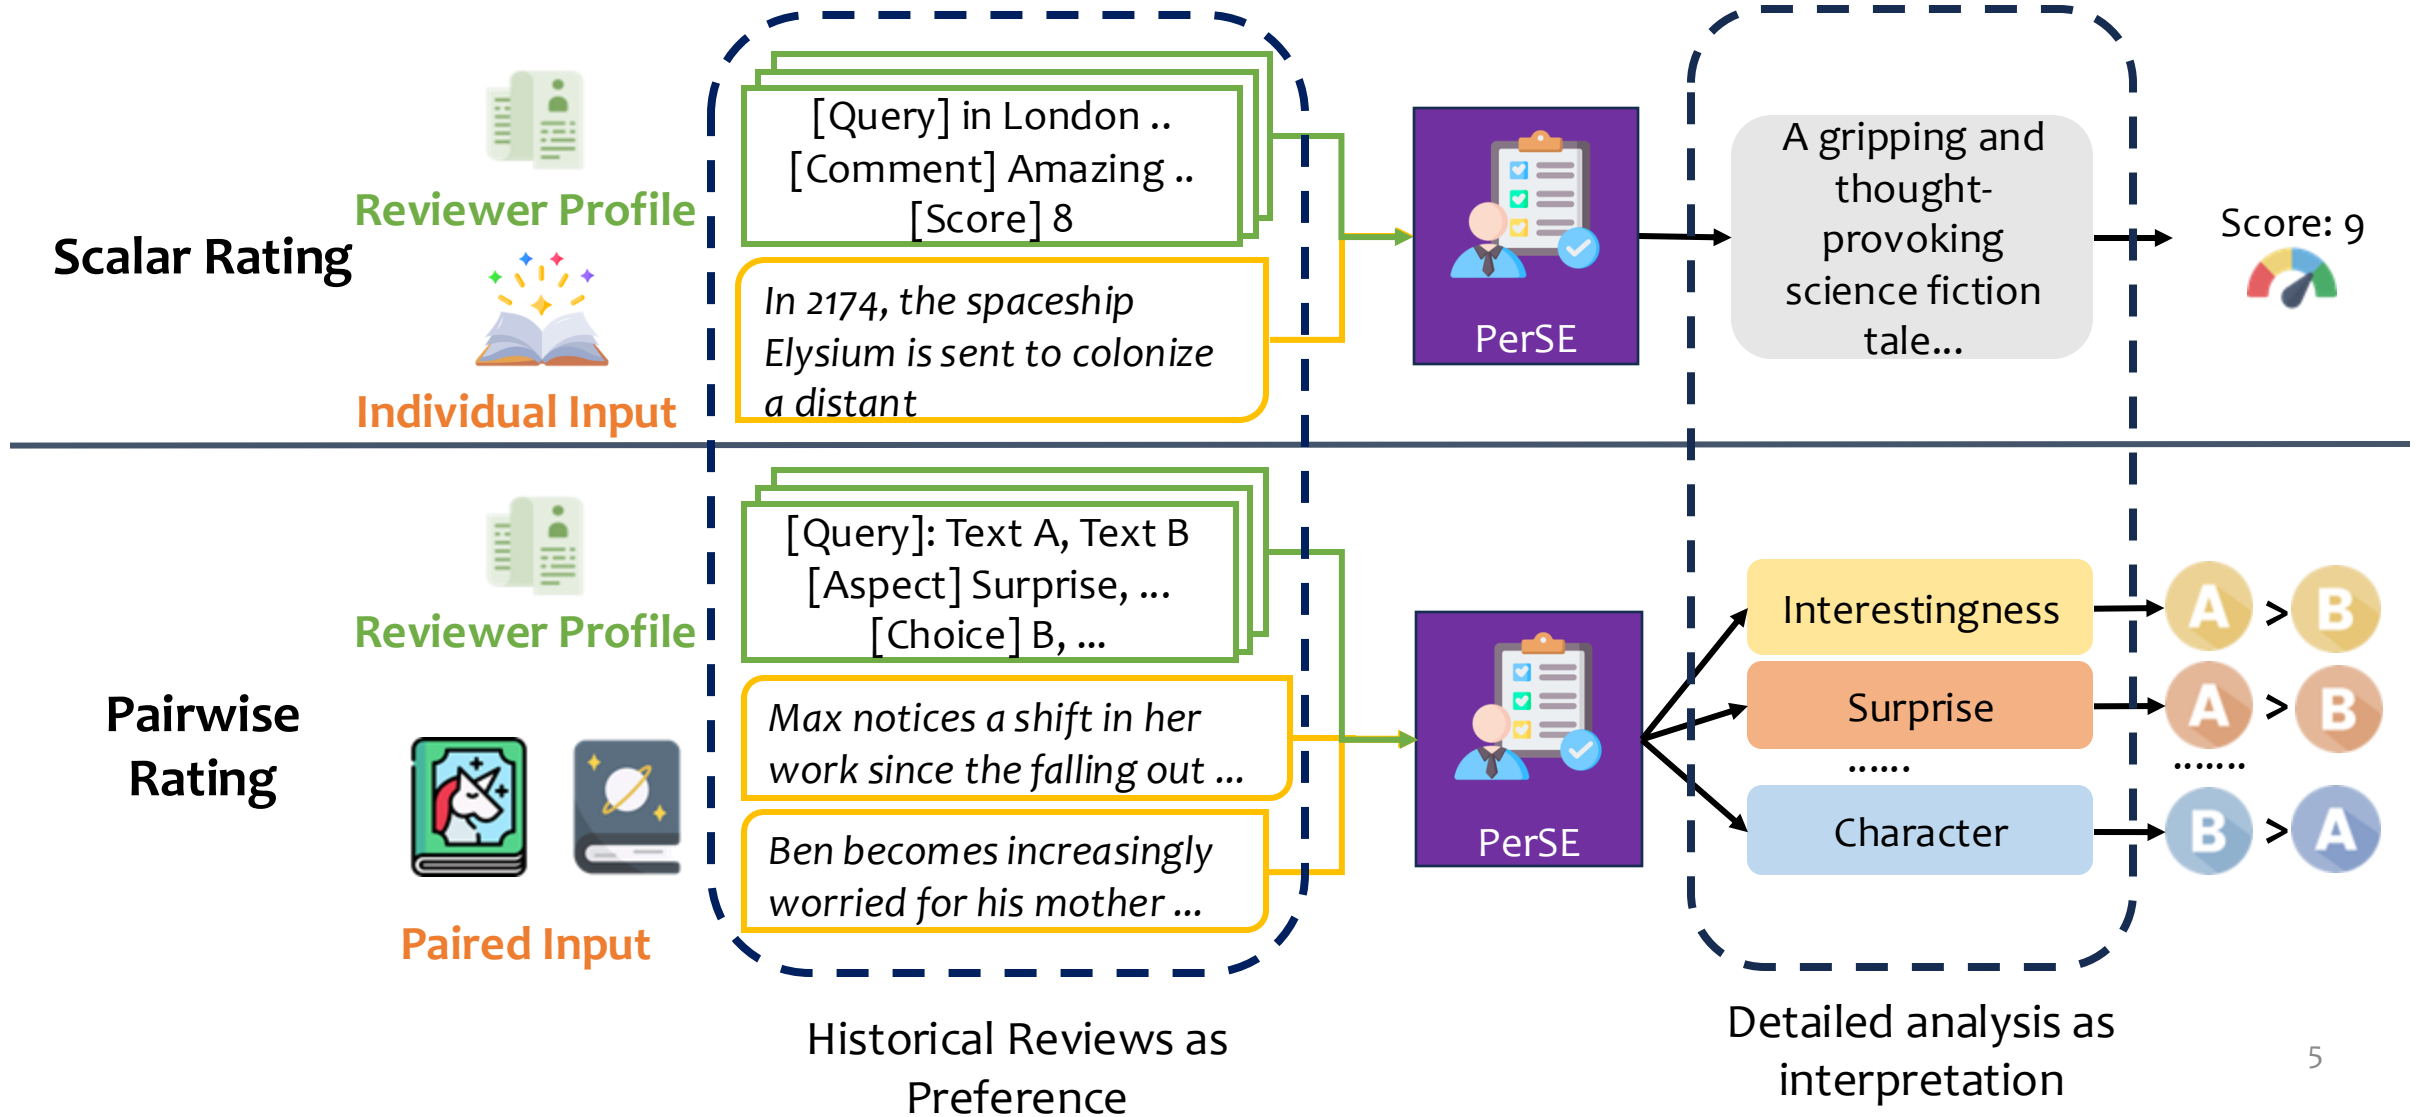
What if we want to tailor generation for personalization?

Learning Personal Preference



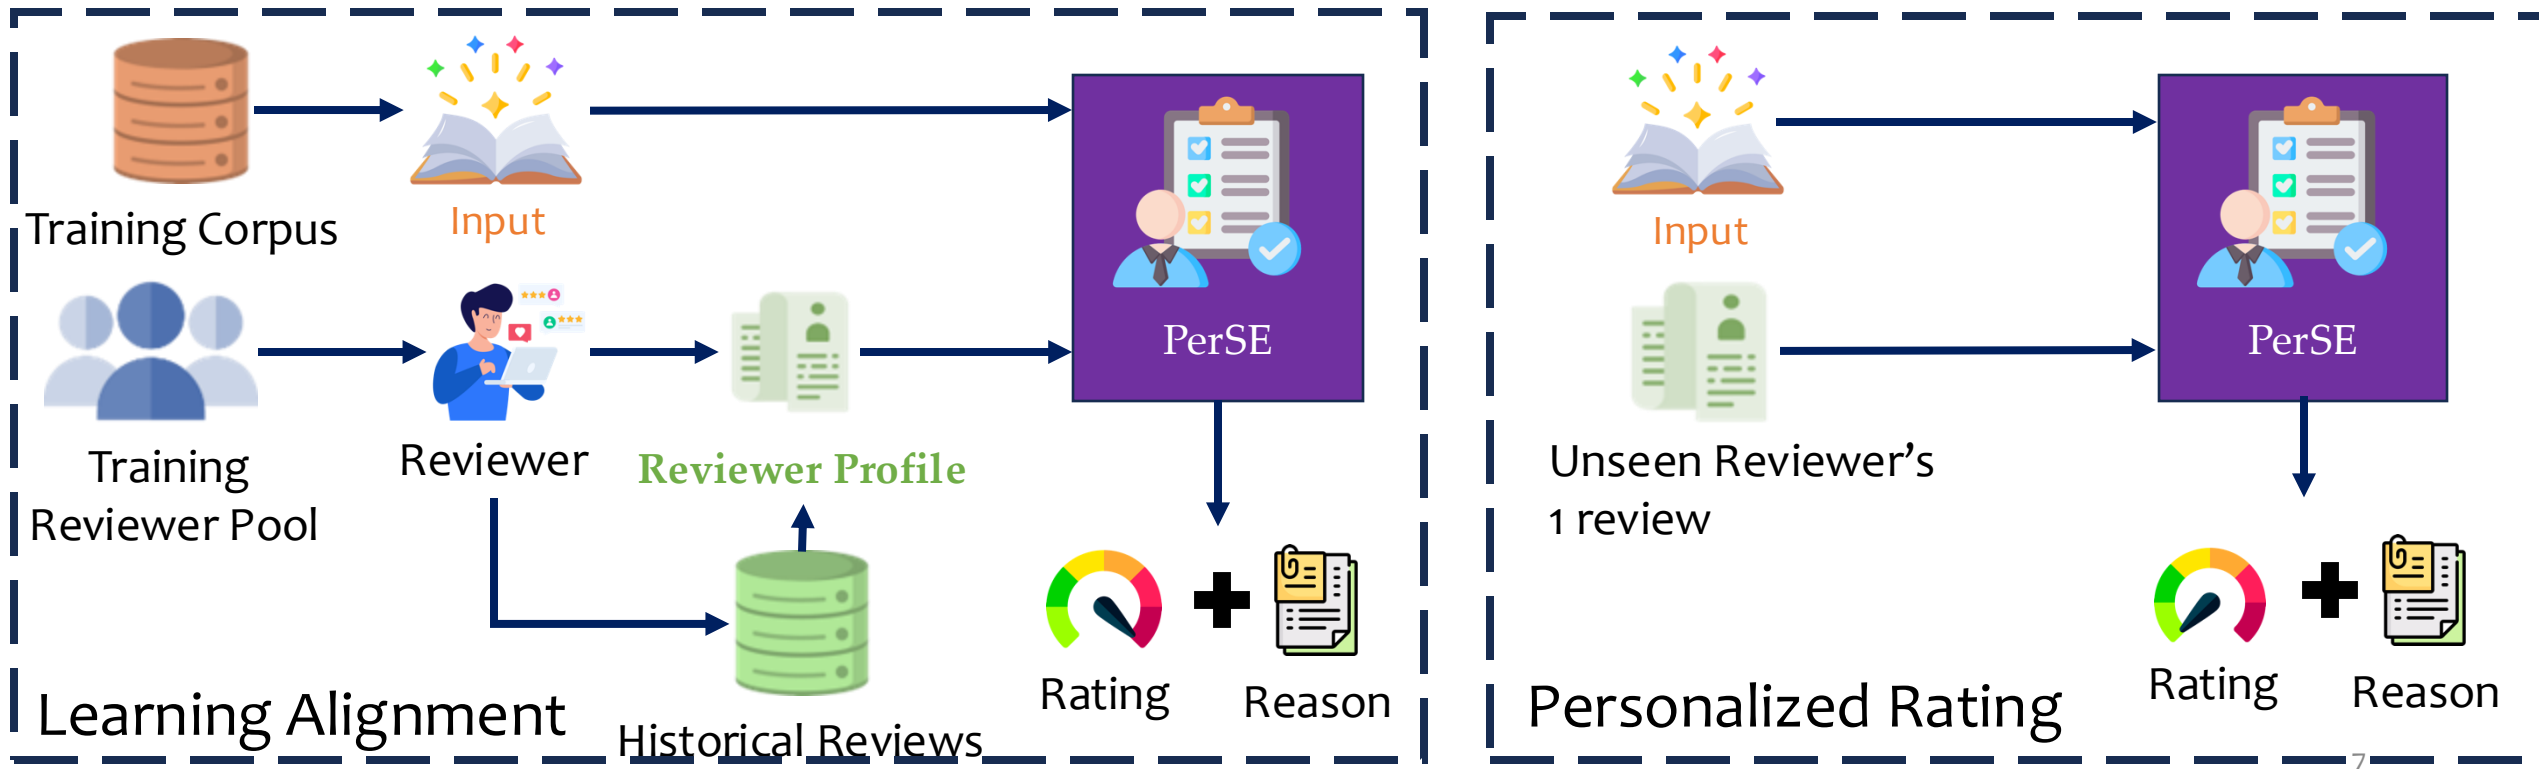
How to collect personal preference from reviewer?
How to infer personalized score from preference?
How to interpret the personalized score?

PerSonalized Evaluation Model (PerSE)




Learning Personalized Alignment in Evaluation

Enhance LLM's capability in aligning with specific preference with limited personal context and detailed explanation




PerSE significantly outperforms all baselines

Highest Correlation with Human Rating on Scalar Rating




Best accuracy for Pairwise Rating on different aspects



PerSE Generalizes Well to Other Domains without Fine-tuning

Directly Apply PerSE trained on Per-MPST to Amazon Book dataset



PerSE cares more about reviewer's preference, while GPT-4 is more general

Reviewer Preference

A newscast is shown about a year after a prehistoric piranha attack on Lake Victoria. The lake is now uninhabitable, and the town has been abandoned due to the drying up of the

critical about the plots, and especially cares about the novelty and vision

you just hate it when this happens? Now that you finally found yourself a gorgeous looking high school sweetheart, she turns out to be a murderous witch, capable of making you eat glass beer bottles with the touch of her hand! [Score] 5

Plots

Professor Jack Hanson and his colleagues go to the Gibraltar ice shelf to take ice core samples and that is poorly recorded. *Big topic but boring* presentation to the UN. The Atlantic show a sharp drop in ocean levels. He believes Jack's theories and works with him to build a forecast model....

Reviewer

I mean, you thing may be to the audience disastrous. *Not appealing, repeated Score: 2* The whole appealing old, gain.

GPT-4

This film offers a chilling, literally and metaphorically, journey into the devastating effects of climate change. *Good topic Score: 7* It's often brushed aside... fear of the unknown future of our planet to craft an engaging narrative. [Score] 7


PerSE

[Review] I movies are. *visible budget restrictions score: 3* the tch, s. Too a

Take Away

- ❖ LLMs finetuned with human feedback is general
 - Learn common interest but struggle with diverse preferences
- ❖ PerSE: instruction tune on a small amount of preference data
 - help LLMs better align with specific preferences

Interested to Learn More?



Paper



Code